

```
In [ ]: import pandas as pd
import numpy as np
import seaborn as sns
import statsmodels.api as sms
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [ ]: df = pd.read_csv('airline.csv')
```

```
In [ ]: df.head()
```

```
Out[ ]:
```

	Month	Thousands of Passengers
0	1949-01	112.0
1	1949-02	118.0
2	1949-03	132.0
3	1949-04	129.0
4	1949-05	121.0

```
In [ ]: df.shape
```

```
Out[ ]: (145, 2)
```

```
In [ ]: df.isnull().sum()
```

```
Out[ ]: Month      0
Thousands of Passengers  1
dtype: int64
```

One missing value in the column Thousands of Passengers. So we will drop it.

```
In [ ]: df.dropna(axis=0, inplace=True)
```

```
In [ ]: df.isnull().sum()
```

```
Out[ ]: Month      0
Thousands of Passengers  0
dtype: int64
```

Now there is no missing value.

Checking the information about the dataset.

```
In [ ]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Index: 144 entries, 0 to 143
Data columns (total 2 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Month                                144 non-null    object
1   Thousands of Passengers              144 non-null    float64
dtypes: float64(1), object(1)
memory usage: 3.4+ KB

```

Month column type is object so we need to convert it into date time.

```
In [ ]: df['Month'] = pd.to_datetime(df['Month'])
```

```
In [ ]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Index: 144 entries, 0 to 143
Data columns (total 2 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Month                                144 non-null    datetime64[ns]
1   Thousands of Passengers              144 non-null    float64
dtypes: datetime64[ns](1), float64(1)
memory usage: 3.4 KB

```

```
In [ ]: df.head()
```

```
Out[ ]:
```

	Month	Thousands of Passengers
0	1949-01-01	112.0
1	1949-02-01	118.0
2	1949-03-01	132.0
3	1949-04-01	129.0
4	1949-05-01	121.0

```
In [ ]: df.set_index('Month',inplace=True)
```

```
In [ ]: df.head()
```


```
Out[ ]:
```

	Thousands of Passengers
Month	
1949-01-01	112.0
1949-02-01	118.0
1949-03-01	132.0
1949-04-01	129.0
1949-05-01	121.0

Let us plot the data

```
In [ ]: df.plot()
```

```
Out[ ]: <Axes: xlabel='Month'>
```



The dataset is non-stationary.

Applying Dickey Fuller test.

```
In [ ]: from statsmodels.tsa.stattools import adfuller
```

```
In [ ]: def adf_test(series):
    result = adfuller(series)
    print('ADF Statistics: {}'.format(result[0]))
    print('p-value: {}'.format(result[1]))

    if result[1] <= 0.05:
        print('Reject the null hypothesis and hence the data is stationary.')
    else:
        print('The null hypothesis is accepted and hence the data is non-stationary')
```

```
In [ ]: adf_test(df['Thousands of Passengers'])
```

ADF Statistics: 0.8153688792060482

p-value: 0.991880243437641

The null hypothesis is accepted and hence the data is non-stationary.

Using the differencing technique

```
In [ ]: df['First Difference'] = df['Thousands of Passengers'] - df['Thousands of Passen
```

```
In [ ]: df.head()
```

```
Out[ ]:
```

	Thousands of Passengers	First Difference
Month		
1949-01-01	112.0	NaN
1949-02-01	118.0	6.0
1949-03-01	132.0	14.0
1949-04-01	129.0	-3.0
1949-05-01	121.0	-8.0

Furhter Applying Dickey-Fuller test

```
In [ ]: adf_test(df['First Difference'].dropna())
```

ADF Statistics: -2.8292668241699994

p-value: 0.0542132902838255

The null hypothesis is accpeted and hence the data is non-stationary.

Still the data is non-stationary so further using differencing technique.

```
In [ ]: df['Second Difference'] = df['First Difference'] - df['First Difference'].shift(1)
```

```
In [ ]: adf_test(df['Second Difference'].dropna())
```

ADF Statistics: -16.384231542468505

p-value: 2.7328918500142407e-29

Reject the null hypothesis and hence the data is stationary.

Since the data is seasonal so we are going to to differencing 12 months because sometime it may happen that ARIMA is not working so we use SARIMA instead.

```
In [ ]: df['12 Difference'] = df['Thousands of Passengers'] - df['Thousands of Passenger'].shift(12)
```

```
In [ ]: adf_test(df['12 Difference'].dropna())
```

ADF Statistics: -3.383020726492481


p-value: 0.011551493085514954

Reject the null hypothesis and hence the data is stationary.


Now let's plot Autocorrelation and Partial correlation graph

```
In [ ]: from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
```



```
In [ ]: acf = plot_acf(df['Second Difference'].dropna())
```



```
In [ ]: pacf = plot_pacf(df['Second Difference'].dropna())
```



```
In [ ]: acf12 = plot_acf(df['12 Difference'].dropna())
         pacf12 = plot_pacf(df['12 Difference'].dropna())
```



Split the data for train and test

```
In [ ]: df
```

Out[]:

	Thousands of Passengers	First Difference	Second Difference	12 Difference
Month				
1949-01-01	112.0	NaN	NaN	NaN
1949-02-01	118.0	6.0	NaN	NaN
1949-03-01	132.0	14.0	8.0	NaN
1949-04-01	129.0	-3.0	-17.0	NaN
1949-05-01	121.0	-8.0	-5.0	NaN
...
1960-08-01	606.0	-16.0	-103.0	47.0
1960-09-01	508.0	-98.0	-82.0	45.0
1960-10-01	461.0	-47.0	51.0	54.0
1960-11-01	390.0	-71.0	-24.0	28.0
1960-12-01	432.0	42.0	113.0	27.0

144 rows × 4 columns

```
In [ ]: from datetime import datetime, timedelta
train_dataset_end = datetime(1955,12,1)
test_data_end = datetime(1960,12,1)
```

```
In [ ]: train_data = df[:train_dataset_end]
test_data = df[train_dataset_end + timedelta(days=1) : test_data_end]
```

Prediction

```
In [ ]: pred_start_date = test_data.index[0]
pre_end_date = test_data.index[-1]
```

```
In [ ]: train_data.shape
```

Out[]: (84, 4)

```
In [ ]: test_data.shape
```

Out[]: (60, 4)

```
In [ ]: train_data
```

Out[]:

	Thousands of Passengers	First Difference	Second Difference	12 Difference
Month				
1949-01-01	112.0	NaN	NaN	NaN
1949-02-01	118.0	6.0	NaN	NaN
1949-03-01	132.0	14.0	8.0	NaN
1949-04-01	129.0	-3.0	-17.0	NaN
1949-05-01	121.0	-8.0	-5.0	NaN
...
1955-08-01	347.0	-17.0	-66.0	54.0
1955-09-01	312.0	-35.0	-18.0	53.0
1955-10-01	274.0	-38.0	-3.0	45.0
1955-11-01	237.0	-37.0	1.0	34.0
1955-12-01	278.0	41.0	78.0	49.0

84 rows × 4 columns

```
In [ ]: test_data
```


Out[]:

Month	Thousands of Passengers	First Difference	Second Difference	12 Difference
1956-01-01	284.0	6.0	-35.0	42.0
1956-02-01	277.0	-7.0	-13.0	44.0
1956-03-01	317.0	40.0	47.0	50.0
1956-04-01	313.0	-4.0	-44.0	44.0
1956-05-01	318.0	5.0	9.0	48.0
1956-06-01	374.0	56.0	51.0	59.0
1956-07-01	413.0	39.0	-17.0	49.0
1956-08-01	405.0	-8.0	-47.0	58.0
1956-09-01	355.0	-50.0	-42.0	43.0
1956-10-01	306.0	-49.0	1.0	32.0
1956-11-01	271.0	-35.0	14.0	34.0
1956-12-01	306.0	35.0	70.0	28.0
1957-01-01	315.0	9.0	-26.0	31.0
1957-02-01	301.0	-14.0	-23.0	24.0
1957-03-01	356.0	55.0	69.0	39.0
1957-04-01	348.0	-8.0	-63.0	35.0
1957-05-01	355.0	7.0	15.0	37.0
1957-06-01	422.0	67.0	60.0	48.0
1957-07-01	465.0	43.0	-24.0	52.0
1957-08-01	467.0	2.0	-41.0	62.0

Month	Thousands of Passengers	First Difference	Second Difference	12 Difference
1957-09-01	404.0	-63.0	-65.0	49.0
1957-10-01	347.0	-57.0	6.0	41.0
1957-11-01	305.0	-42.0	15.0	34.0
1957-12-01	336.0	31.0	73.0	30.0
1958-01-01	340.0	4.0	-27.0	25.0
1958-02-01	318.0	-22.0	-26.0	17.0
1958-03-01	362.0	44.0	66.0	6.0
1958-04-01	348.0	-14.0	-58.0	0.0
1958-05-01	363.0	15.0	29.0	8.0
1958-06-01	435.0	72.0	57.0	13.0
1958-07-01	491.0	56.0	-16.0	26.0
1958-08-01	505.0	14.0	-42.0	38.0
1958-09-01	404.0	-101.0	-115.0	0.0
1958-10-01	359.0	-45.0	56.0	12.0
1958-11-01	310.0	-49.0	-4.0	5.0
1958-12-01	337.0	27.0	76.0	1.0
1959-01-01	360.0	23.0	-4.0	20.0
1959-02-01	342.0	-18.0	-41.0	24.0
1959-03-01	406.0	64.0	82.0	44.0
1959-04-01	396.0	-10.0	-74.0	48.0

Month	Thousands of Passengers	First Difference	Second Difference	12 Difference
1959-05-01	420.0	24.0	34.0	57.0
1959-06-01	472.0	52.0	28.0	37.0
1959-07-01	548.0	76.0	24.0	57.0
1959-08-01	559.0	11.0	-65.0	54.0
1959-09-01	463.0	-96.0	-107.0	59.0
1959-10-01	407.0	-56.0	40.0	48.0
1959-11-01	362.0	-45.0	11.0	52.0
1959-12-01	405.0	43.0	88.0	68.0
1960-01-01	417.0	12.0	-31.0	57.0
1960-02-01	391.0	-26.0	-38.0	49.0
1960-03-01	419.0	28.0	54.0	13.0
1960-04-01	461.0	42.0	14.0	65.0
1960-05-01	472.0	11.0	-31.0	52.0
1960-06-01	535.0	63.0	52.0	63.0
1960-07-01	622.0	87.0	24.0	74.0
1960-08-01	606.0	-16.0	-103.0	47.0
1960-09-01	508.0	-98.0	-82.0	45.0
1960-10-01	461.0	-47.0	51.0	54.0
1960-11-01	390.0	-71.0	-24.0	28.0
1960-12-01	432.0	42.0	113.0	27.0

Create an ARIMA model

```
In [ ]: from statsmodels.tsa.arima.model import ARIMA
```

```
In [ ]: arima = ARIMA(train_data['Thousands of Passengers'], order=(0,2,0))
```

```
c:\Users\User\AppData\Local\Programs\Python\Python312\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
  self._init_dates(dates, freq)
c:\Users\User\AppData\Local\Programs\Python\Python312\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
  self._init_dates(dates, freq)
c:\Users\User\AppData\Local\Programs\Python\Python312\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.
  self._init_dates(dates, freq)
```

Fit the model

```
In [ ]: arima_fit = arima.fit()
```

Summary of the Model

```
In [ ]: arima_fit.summary()
```

Out[]: SARIMAX Results

Dep. Variable:	Thousands of Passengers	No. Observations:	84
Model:	ARIMA(0, 2, 0)	Log Likelihood	-385.792
Date:	Sun, 04 Aug 2024	AIC	773.584
Time:	16:00:33	BIC	775.991
Sample:	01-01-1949	HQIC	774.550
	- 12-01-1955		
Covariance Type:	opg		
	coef	std err	z P> z [0.025 0.975]
sigma2	714.5859	102.414	6.977 0.000 513.858 915.314
Ljung-Box (L1) (Q):	4.59	Jarque-Bera (JB):	1.74
Prob(Q):	0.03	Prob(JB):	0.42
Heteroskedasticity (H):	3.19	Skew:	0.31
Prob(H) (two-sided):	0.00	Kurtosis:	3.36

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [ ]: pred_start_date = test_data.index[0]
pre_end_date = test_data.index[-1]
print('The start date is :',pred_start_date)
print('The end date is :',pre_end_date)
```

The start date is : 1956-01-01 00:00:00
The end date is : 1960-12-01 00:00:00

Make Prediction

```
In [ ]: pred = arima_fit.predict(start=pred_start_date, end=pre_end_date)

In [ ]: pred
```

```
Out[ ]: 1956-01-01    319.0
        1956-02-01    360.0
        1956-03-01    401.0
        1956-04-01    442.0
        1956-05-01    483.0
        1956-06-01    524.0
        1956-07-01    565.0
        1956-08-01    606.0
        1956-09-01    647.0
        1956-10-01    688.0
        1956-11-01    729.0
        1956-12-01    770.0
        1957-01-01    811.0
        1957-02-01    852.0
        1957-03-01    893.0
        1957-04-01    934.0
        1957-05-01    975.0
        1957-06-01   1016.0
        1957-07-01   1057.0
        1957-08-01   1098.0
        1957-09-01   1139.0
        1957-10-01   1180.0
        1957-11-01   1221.0
        1957-12-01   1262.0
        1958-01-01   1303.0
        1958-02-01   1344.0
        1958-03-01   1385.0
        1958-04-01   1426.0
        1958-05-01   1467.0
        1958-06-01   1508.0
        1958-07-01   1549.0
        1958-08-01   1590.0
        1958-09-01   1631.0
        1958-10-01   1672.0
        1958-11-01   1713.0
        1958-12-01   1754.0
        1959-01-01   1795.0
        1959-02-01   1836.0
        1959-03-01   1877.0
        1959-04-01   1918.0
        1959-05-01   1959.0
        1959-06-01   2000.0
        1959-07-01   2041.0
        1959-08-01   2082.0
        1959-09-01   2123.0
        1959-10-01   2164.0
        1959-11-01   2205.0
        1959-12-01   2246.0
        1960-01-01   2287.0
        1960-02-01   2328.0
        1960-03-01   2369.0
        1960-04-01   2410.0
        1960-05-01   2451.0
        1960-06-01   2492.0
        1960-07-01   2533.0
        1960-08-01   2574.0
        1960-09-01   2615.0
        1960-10-01   2656.0
        1960-11-01   2697.0
```

1960-12-01 2738.0

Freq: MS, Name: predicted_mean, dtype: float64

Residuals

```
In [ ]: residuals=test_data['Thousands of Passengers']-pred
```

```
In [ ]: residuals
```


```
Out[ ]: Month
1956-01-01      -35.0
1956-02-01      -83.0
1956-03-01      -84.0
1956-04-01     -129.0
1956-05-01     -165.0
1956-06-01     -150.0
1956-07-01     -152.0
1956-08-01     -201.0
1956-09-01     -292.0
1956-10-01     -382.0
1956-11-01     -458.0
1956-12-01     -464.0
1957-01-01     -496.0
1957-02-01     -551.0
1957-03-01     -537.0
1957-04-01     -586.0
1957-05-01     -620.0
1957-06-01     -594.0
1957-07-01     -592.0
1957-08-01     -631.0
1957-09-01     -735.0
1957-10-01     -833.0
1957-11-01     -916.0
1957-12-01     -926.0
1958-01-01     -963.0
1958-02-01    -1026.0
1958-03-01    -1023.0
1958-04-01    -1078.0
1958-05-01    -1104.0
1958-06-01    -1073.0
1958-07-01    -1058.0
1958-08-01    -1085.0
1958-09-01    -1227.0
1958-10-01    -1313.0
1958-11-01    -1403.0
1958-12-01    -1417.0
1959-01-01    -1435.0
1959-02-01    -1494.0
1959-03-01    -1471.0
1959-04-01    -1522.0
1959-05-01    -1539.0
1959-06-01    -1528.0
1959-07-01    -1493.0
1959-08-01    -1523.0
1959-09-01    -1660.0
1959-10-01    -1757.0
1959-11-01    -1843.0
1959-12-01    -1841.0
1960-01-01    -1870.0
1960-02-01    -1937.0
1960-03-01    -1950.0
1960-04-01    -1949.0
1960-05-01    -1979.0
1960-06-01    -1957.0
1960-07-01    -1911.0
1960-08-01    -1968.0
1960-09-01    -2107.0
1960-10-01    -2195.0
1960-11-01    -2307.0
```



```
1960-12-01    -2306.0  
dtype: float64
```

```
In [ ]: arima_fit.resid.plot(kind='kde')
```

```
Out[ ]: <Axes: ylabel='Density'>
```



The plot of residuals follows normal distribution

```
In [ ]: test_data['Predicted_ARIMA'] = pred
```

```
C:\Users\User\AppData\Local\Temp\ipykernel_18036\284031954.py:1: SettingWithCopyWarning:  
  A value is trying to be set on a copy of a slice from a DataFrame.  
  Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy  
test_data['Predicted_ARIMA'] = pred
```

```
In [ ]: test_data.head()
```


Out[]:

	Thousands of Passengers	First Difference	Second Difference	12 Difference	Predicted_ARIMA
Month					
1956-01-01	284.0	6.0	-35.0	42.0	319.0
1956-02-01	277.0	-7.0	-13.0	44.0	360.0
1956-03-01	317.0	40.0	47.0	50.0	401.0
1956-04-01	313.0	-4.0	-44.0	44.0	442.0
1956-05-01	318.0	5.0	9.0	48.0	483.0

Let us plot Thousands of Passengers vs Predicted_ARIMA

```
In [ ]: test_data[['Thousands of Passengers', 'Predicted_ARIMA']].plot()
```

Out[]: <Axes: xlabel='Month'>



This seems very poor prediction.

Let us create a SARIMAX Model now.

```
In [ ]: from statsmodels.tsa.statespace.sarimax import SARIMAX
```

```
In [ ]: sarima = SARIMAX(train_data['Thousands of Passengers'], order=(3,0,5), seasonal_
```

```
c:\Users\User\AppData\Local\Programs\Python\Python312\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.  
    self._init_dates(dates, freq)  
c:\Users\User\AppData\Local\Programs\Python\Python312\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency MS will be used.  
    self._init_dates(dates, freq)
```

Fit the model

```
In [ ]: sarima_fit = sarima.fit()
```

```
c:\Users\User\AppData\Local\Programs\Python\Python312\Lib\site-packages\statsmodels\tsa\base\model.py:607: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals  
    warnings.warn("Maximum Likelihood optimization failed to "
```

Summary

```
In [ ]: sarima_fit.summary()
```

Out[]:

SARIMAX Results

Dep. Variable:		Thousands of Passengers			No. Observations:		84
Model:		SARIMAX(3, 0, 5)x(0, 1, [], 12)			Log Likelihood		-265.240
Date:		Sun, 04 Aug 2024			AIC		548.481
Time:		16:00:33			BIC		568.971
Sample:		01-01-1949			HQIC		556.638
		- 12-01-1955					
Covariance Type:		opg					
	coef	std err	z	P> z 	[0.025	0.975]	
ar.L1	0.5983	0.937	0.638	0.523	-1.239	2.436	
ar.L2	0.8311	0.232	3.581	0.000	0.376	1.286	
ar.L3	-0.4525	0.894	-0.506	0.613	-2.204	1.299	
ma.L1	0.1837	1.165	0.158	0.875	-2.099	2.467	
ma.L2	-0.5341	1.263	-0.423	0.672	-3.009	1.940	
ma.L3	-0.0986	0.384	-0.257	0.798	-0.852	0.655	
ma.L4	-0.1273	0.338	-0.377	0.706	-0.789	0.535	
ma.L5	0.2471	0.357	0.693	0.489	-0.452	0.947	
sigma2	87.7323	81.217	1.080	0.280	-71.451	246.915	
Ljung-Box (L1) (Q):		0.02	Jarque-Bera (JB):		2.68		
Prob(Q):		0.88	Prob(JB):		0.26		
Heteroskedasticity (H):		2.05	Skew:		0.46		
Prob(H) (two-sided):		0.09	Kurtosis:		2.77		

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

In []: test_data.head()

Out[]:	Thousands of Passengers	First Difference	Second Difference	12 Difference	Predicted_ARIMA
Month					
1956-01-01	284.0	6.0	-35.0	42.0	319.0
1956-02-01	277.0	-7.0	-13.0	44.0	360.0
1956-03-01	317.0	40.0	47.0	50.0	401.0
1956-04-01	313.0	-4.0	-44.0	44.0	442.0
1956-05-01	318.0	5.0	9.0	48.0	483.0

Prediction

```
In [ ]: pred_start_date = test_data.index[0]
pre_end_date = test_data.index[-1]
print('The start date is :',pred_start_date)
print('The end date is :',pre_end_date)
```

The start date is : 1956-01-01 00:00:00
The end date is : 1960-12-01 00:00:00


```
In [ ]: pred_Sarima=sarima_fit.predict(start=datetime(1956,6,6),end=datetime(1960,12,1))
```

Residuals

```
In [ ]: residuals=test_data['Thousands of Passengers']-pred_Sarima
```


```
In [ ]: sarima_fit.resid.plot()
```

Out[]: <Axes: xlabel='Month'>



```
In [ ]: sarima_fit.resid.plot(kind='kde')
```

```
Out[ ]: <Axes: ylabel='Density'>
```



```
In [ ]: test_data['Predicted_SARIMA']=pred_Sarima
```


```
C:\Users\User\AppData\Local\Temp\ipykernel_18036\1367177785.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
test_data['Predicted_SARIMA']=pred_Sarima
```

```
In [ ]: test_data.head()
```

```
Out[ ]:
```

	Thousands of Passengers	First Difference	Second Difference	12 Difference	Predicted_ARIMA	Predicted_SARIMA
Month						
1956-01-01	284.0	6.0	-35.0	42.0	319.0	
1956-02-01	277.0	-7.0	-13.0	44.0	360.0	
1956-03-01	317.0	40.0	47.0	50.0	401.0	
1956-04-01	313.0	-4.0	-44.0	44.0	442.0	
1956-05-01	318.0	5.0	9.0	48.0	483.0	



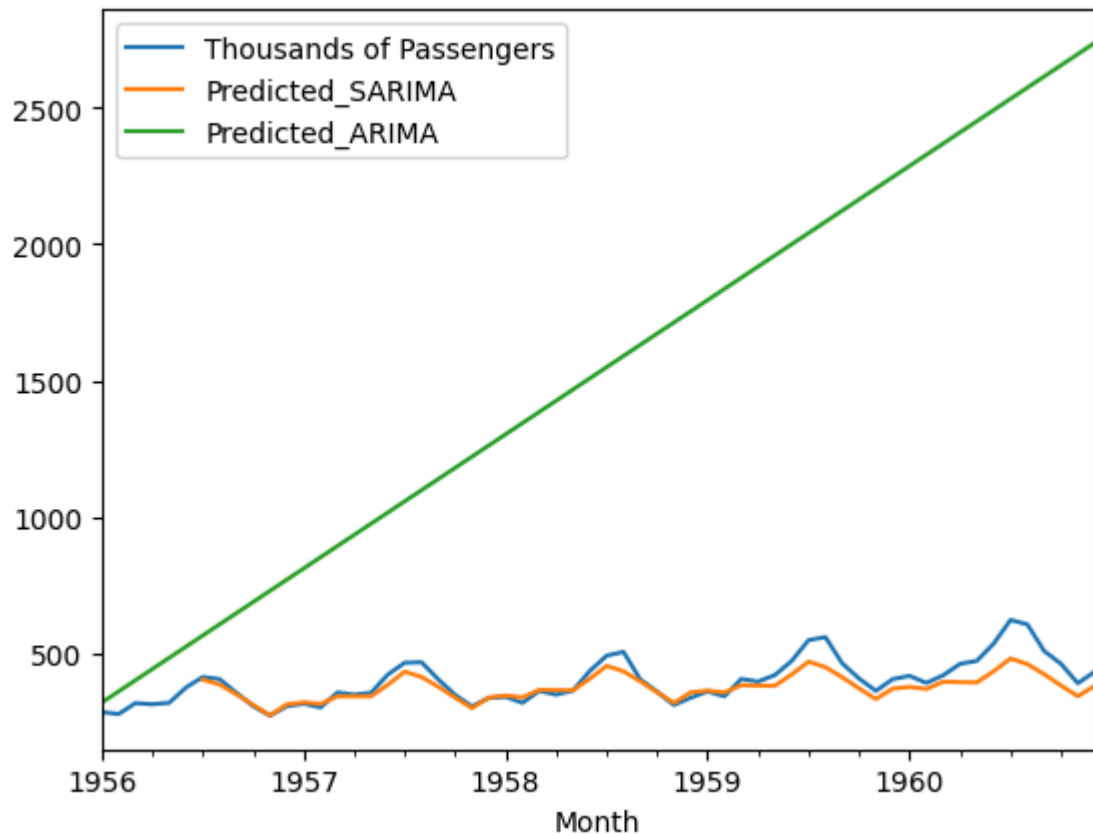
```
In [ ]: test_data.shape
```

```
Out[ ]: (60, 6)
```

Let us plot Thousands of Passengers vs Predicted_ARIMA vs Predicted_SARIMA

```
In [ ]: test_data[['Thousands of Passengers', 'Predicted_SARIMA', 'Predicted_ARIMA']].plot
```

```
Out[ ]: <Axes: xlabel='Month'>
```



SARIMAX model fits well to the test data.

Let us try with Random Forest

```
In [ ]: df1 = pd.read_csv('airline.csv')
```

```
In [ ]: df1.head()
```

```
Out[ ]:
```

	Month	Thousands of Passengers
0	1949-01	112.0
1	1949-02	118.0
2	1949-03	132.0
3	1949-04	129.0
4	1949-05	121.0

```
In [ ]: df1.shape
```

```
Out[ ]: (145, 2)
```

```
In [ ]: df1.isnull().sum()
```

```
Out[ ]: Month          0
Thousands of Passengers  1
dtype: int64
```

```
In [ ]: df1.dropna(axis=0, inplace=True)
```



```
In [ ]: df1.shape
```

```
Out[ ]: (144, 2)
```

```
In [ ]: df1['Month'] = pd.to_datetime(df1['Month'])
df1.set_index('Month', inplace=True)
print(df1.head())
```

	Thousands of Passengers
Month	
1949-01-01	112.0
1949-02-01	118.0
1949-03-01	132.0
1949-04-01	129.0
1949-05-01	121.0

```
In [ ]: df1.columns
```

```
Out[ ]: Index(['Thousands of Passengers'], dtype='object')
```

Let us create features

```
In [ ]: def create_features(df1):
    df1['Month'] = df1.index.month
    df1['Year'] = df1.index.year
    for i in range(1, 13):
        df1[f'Lag_{i}'] = df1['Thousands of Passengers'].shift(i)
    df1.dropna(inplace=True)
    return df1

df1 = create_features(df1)
print(df1.head())
```

	Thousands of Passengers	Month	Year	Lag_1	Lag_2	Lag_3	Lag_4	\
Month								
1950-01-01	115.0	1	1950	118.0	104.0	119.0	136.0	
1950-02-01	126.0	2	1950	115.0	118.0	104.0	119.0	
1950-03-01	141.0	3	1950	126.0	115.0	118.0	104.0	
1950-04-01	135.0	4	1950	141.0	126.0	115.0	118.0	
1950-05-01	125.0	5	1950	135.0	141.0	126.0	115.0	

	Lag_5	Lag_6	Lag_7	Lag_8	Lag_9	Lag_10	Lag_11	Lag_12
Month								
1950-01-01	148.0	148.0	135.0	121.0	129.0	132.0	118.0	112.0
1950-02-01	136.0	148.0	148.0	135.0	121.0	129.0	132.0	118.0
1950-03-01	119.0	136.0	148.0	148.0	135.0	121.0	129.0	132.0
1950-04-01	104.0	119.0	136.0	148.0	148.0	135.0	121.0	129.0
1950-05-01	118.0	104.0	119.0	136.0	148.0	148.0	135.0	121.0

Train Test Split

```
In [ ]: from sklearn.model_selection import train_test_split
```

```
In [ ]: X = df1.drop(columns='Thousands of Passengers', axis=1)
y = df1['Thousands of Passengers']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, shuffle
```

```
In [ ]: X.shape
```

```
Out[ ]: (132, 14)
```

```
In [ ]: X_train.shape
```

```
Out[ ]: (105, 14)
```

```
In [ ]: y.shape
```

```
Out[ ]: (132,)
```

```
In [ ]: X.head()
```

```
Out[ ]:
```

	Month	Year	Lag_1	Lag_2	Lag_3	Lag_4	Lag_5	Lag_6	Lag_7	Lag_8	Lag_9
Month											
1950-01-01	1	1950	118.0	104.0	119.0	136.0	148.0	148.0	135.0	121.0	129.0
1950-02-01	2	1950	115.0	118.0	104.0	119.0	136.0	148.0	148.0	135.0	121.0
1950-03-01	3	1950	126.0	115.0	118.0	104.0	119.0	136.0	148.0	148.0	135.0
1950-04-01	4	1950	141.0	126.0	115.0	118.0	104.0	119.0	136.0	148.0	148.0
1950-05-01	5	1950	135.0	141.0	126.0	115.0	118.0	104.0	119.0	136.0	148.0

◀ ————— ▶

```
In [ ]: y.head()
```

```
Out[ ]: Month
1950-01-01    115.0
1950-02-01    126.0
1950-03-01    141.0
1950-04-01    135.0
1950-05-01    125.0
Name: Thousands of Passengers, dtype: float64
```

```
In [ ]: from sklearn.ensemble import RandomForestRegressor
```

```
In [ ]: rf = RandomForestRegressor(n_estimators=100, random_state=42)
```

```
In [ ]: rf.fit(X_train, y_train)
```

```
Out[ ]: ▼ RandomForestRegressor
RandomForestRegressor(random_state=42)
```

```
In [ ]: rf_pred = rf.predict(X_test)
```

```
In [ ]: from sklearn.metrics import mean_squared_error, mean_absolute_percentage_error
```


```
In [ ]: print('MSE :', mean_squared_error(y_test,rf_pred))  
print('MAPE :', mean_absolute_percentage_error(y_test,rf_pred))
```

MSE : 2274.1612666666666

MAPE : 0.07027829414813154

```
In [ ]: test_data[['Thousands of Passengers','Predicted_SARIMA','Predicted_ARIMA']].plot
```

Out[]: <Axes: xlabel='Month'>



```
In [ ]:
```